



US007555142B2

(12) **United States Patent**
Hildreth et al.

(10) **Patent No.:** **US 7,555,142 B2**
(45) **Date of Patent:** ***Jun. 30, 2009**

(54) **MULTIPLE CAMERA CONTROL SYSTEM**

(75) Inventors: **Evan Hildreth**, Ottawa (CA); **Francis MacDougall**, Ottawa (CA)

(73) Assignee: **GestureTek, Inc.**, Sunnyvale, CA (US)

(*) Notice: Subject to any disclaimer, the term of this patent is extended or adjusted under 35 U.S.C. 154(b) by 0 days.

This patent is subject to a terminal disclaimer.

(21) Appl. No.: **11/932,869**

(22) Filed: **Oct. 31, 2007**

(65) **Prior Publication Data**

US 2008/0056536 A1 Mar. 6, 2008

Related U.S. Application Data

(63) Continuation of application No. 11/304,000, filed on Dec. 19, 2005, now Pat. No. 7,421,093, which is a continuation of application No. 09/962,612, filed on Sep. 26, 2001, now Pat. No. 7,058,204.

(60) Provisional application No. 60/237,187, filed on Oct. 3, 2000.

(51) **Int. Cl.**

G06K 9/00 (2006.01)

G06K 9/36 (2006.01)

(52) **U.S. Cl.** **382/103; 382/154**

(58) **Field of Classification Search** **382/103, 382/106, 107, 154, 286, 291, 292; 348/42, 348/47, 48, 143, 153, 159; 345/863**

See application file for complete search history.

(56) **References Cited**

U.S. PATENT DOCUMENTS

4,144,449 A 3/1979 Funk et al.

(Continued)

FOREIGN PATENT DOCUMENTS

DE 198 10 452 A1 12/1998

(Continued)

OTHER PUBLICATIONS

Z. Zhang; "Determining the Epipolar Geometry and its Uncertainty"; A Review, The International Journal of Computer Vision, 1997, INRIA Sophia Antipolis, France, Jul. 1996, pp. 161-195.

(Continued)


Primary Examiner—Yon Couso

(74) *Attorney, Agent, or Firm*—Fish & Richardson P.C.

(57) **ABSTRACT**

A multiple camera tracking system for interfacing with an application program running on a computer is provided. The tracking system includes two or more video cameras arranged to provide different viewpoints of a region of interest, and are operable to produce a series of video images. A processor is operable to receive the series of video images and detect objects appearing in the region of interest. The processor executes a process to generate a background data set from the video images, generate an image data set for each received video image, compare each image data set to the background-data set to produce a difference map for each image data set, detect a relative position of an object of interest within each difference map, and produce an absolute position of the object of interest from the relative positions of the object of interest and map the absolute position to a position indicator associated with the application program.

35 Claims, 15 Drawing Sheets



U.S. PATENT DOCUMENTS

4,247,767 A 1/1981 O'Brien et al.
 4,507,557 A 3/1985 Tsikos
 4,742,221 A 5/1988 Sasaki et al.
 4,746,770 A 5/1988 McAvinney
 4,782,328 A 11/1988 Denlinger
 4,818,826 A 4/1989 Kimura
 5,097,516 A 3/1992 Amir
 5,109,435 A 4/1992 Lo et al.
 5,317,140 A 5/1994 Dunthorn
 5,483,603 A 1/1996 Luke et al.
 5,502,568 A 3/1996 Ogawa et al.
 5,534,917 A 7/1998 MacDougall
 5,581,276 A 12/1996 Cipolla et al.
 5,638,092 A 6/1997 Eng et al.
 5,769,640 A 6/1998 Jacobus et al.
 5,844,392 A 12/1998 Peurach et al.
 5,907,328 A 5/1999 Brush, II et al.
 5,911,004 A 6/1999 Chuchi et al.
 5,936,615 A 8/1999 Waters
 6,002,808 A 12/1999 Freeman
 6,008,798 A 12/1999 Mato, Jr. et al.
 6,072,494 A 6/2000 Nguyen
 6,075,895 A 6/2000 Qiao et al.
 6,100,538 A 8/2000 Ogawa
 6,173,066 B1 1/2001 Peurach et al.
 6,191,773 B1 2/2001 Maruno et al.
 6,208,330 B1 3/2001 Hasegawa et al.
 6,215,890 B1 4/2001 Matsuo et al.
 6,222,465 B1 4/2001 Kumar et al.
 6,256,033 B1 7/2001 Nguyen
 6,335,724 B1 1/2002 Takekawa et al.
 6,414,671 B1 7/2002 Gillespie et al.
 6,421,042 B1 7/2002 Omura et al.
 6,429,856 B1 8/2002 Omura et al.
 6,531,999 B1 3/2003 Trajkovic
 6,563,491 B1 5/2003 Omura
 6,594,023 B1 7/2003 Omura et al.
 6,993,179 B1 1/2006 Weinshall et al.
 2001/0019325 A1 9/2001 Takekawa
 2001/0020933 A1 9/2001 Maggioni
 2001/0022579 A1 9/2001 Hirabayashi
 2001/0026268 A1 10/2001 Ito
 2002/0015024 A1 2/2002 Westerman et al.
 2003/0112228 A1 6/2003 Gillespie et al.
 2004/0012573 A1 1/2004 Morrison et al.

FOREIGN PATENT DOCUMENTS

EP 0 905 644 A2 3/1999

JP 57-211637 12/1982
 JP 8-240407 9/1996
 JP 9-91094 4/1997
 JP 9-319501 12/1997
 WO WO 99/35633 7/1999
 WO WO 99/40562 8/1999

OTHER PUBLICATIONS

C. Urmson; "A Comparison of Pt. Grey Research's Digiclops and Videre Design's Small Vision System for Sensing on the Hyperion Robot"; ;Nov. 2000 (<http://www.videredesign.com/StereoComparison.pdf>).
 M. Okutomi and T. Kanda. "A Multiple-baseline Stereo", PAMI 15(4), Apr. 1993, pp. 355-363.
 D. Scharstein and R. Szeliski, "Stereo Matching with Non-Linear Diffusion" which was published in Proceedings of the 1996 Conference on Computer Vision and Pattern Recognition (CVPR '96) middlebury.edu/~schar/papers/diffusion.ps.gz.
 G. Morrison et al. "Machine Vision Passive Touch Technology for Interactive Displays" SID 01 Digest; Jun. 2001, pp. 74-77.
 XP-002229564—"Robust Finger Tracking with Multiple Cameras"; Cullen Jennings, University of British Columbia, Vancouver, B.C., Canada; 1999; pp. 152-160.
 XP-002229565—"Wallflower: Principles and Practice of Background Maintenance"; Kentaro Toyama et al.; Microsoft Research; Redmond, WA; 1999; vol. 1 —pp. 255-261.
 XP-002229566—"Background Modeling for Segmentation of Video-Rate Stereo Sequences"; Christopher Eveland, et al.; Computer Science Department, University of Rochester, Rochester, NY; 1998; pp. 266-271.
 Bud K. Funk, *CCDs in Optical Touch Panels Deliver High Resolution*; Sep. 1980; Electronic Design.
 Z. Zhang; "A Flexible New Technique for Camera Calibration"; *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 22, No. 11, pp. 1330-1334, Dec. 2, 1998, Microsoft Research; <http://research.microsoft.com/~zhang>.
 Point Grey Research Inc. Digiclops (<http://www.ptgrey.com/>), 2000, PGR, Inc.
 Small Vision System Stereo Developer Kit (<http://www.videredesign.com>), Feb. 2, 2001.
 C. Jennings. "Robust Finger Tracking with Multiple Cameras" Proceeding, International Workshop on Recognition, Analysis, and Tracking of Faces and Gestures in Real-Time Systems, Jul. 20, 1999, (<http://www.cs.ubc.ca/spider/jennings/ratfg-rts99/cj99.html>).
 XP-10378301—"Detection and Estimation of Pointing Gestures in Dense Disparity Maps"; Nebojsa Jojic et al.; Fourth IEEE International Conference on Automatic Face and Gesture Recognition, Mar. 28-30, 2000, ECE Dept. & Beckman Institute; University of Illinois, Urbana, Ill; pp. 468-475.




Fig. 1




Fig. 2A




Fig. 2B

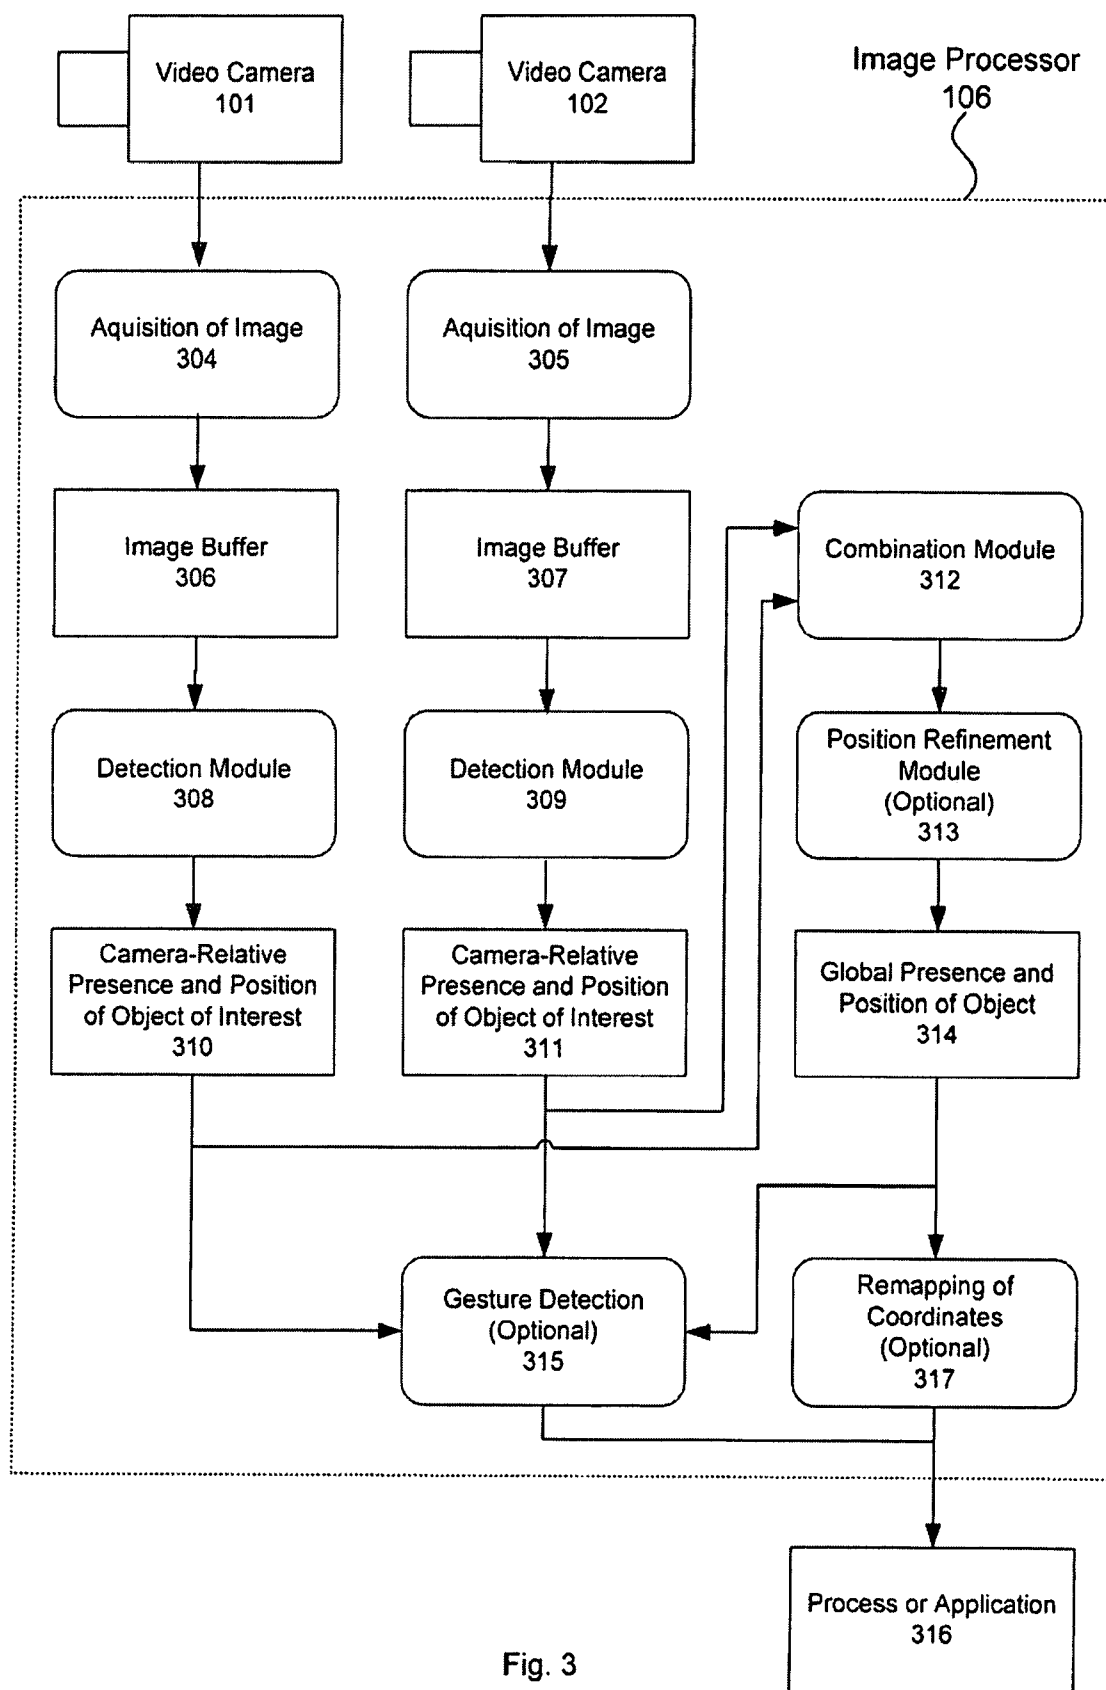


Fig. 3

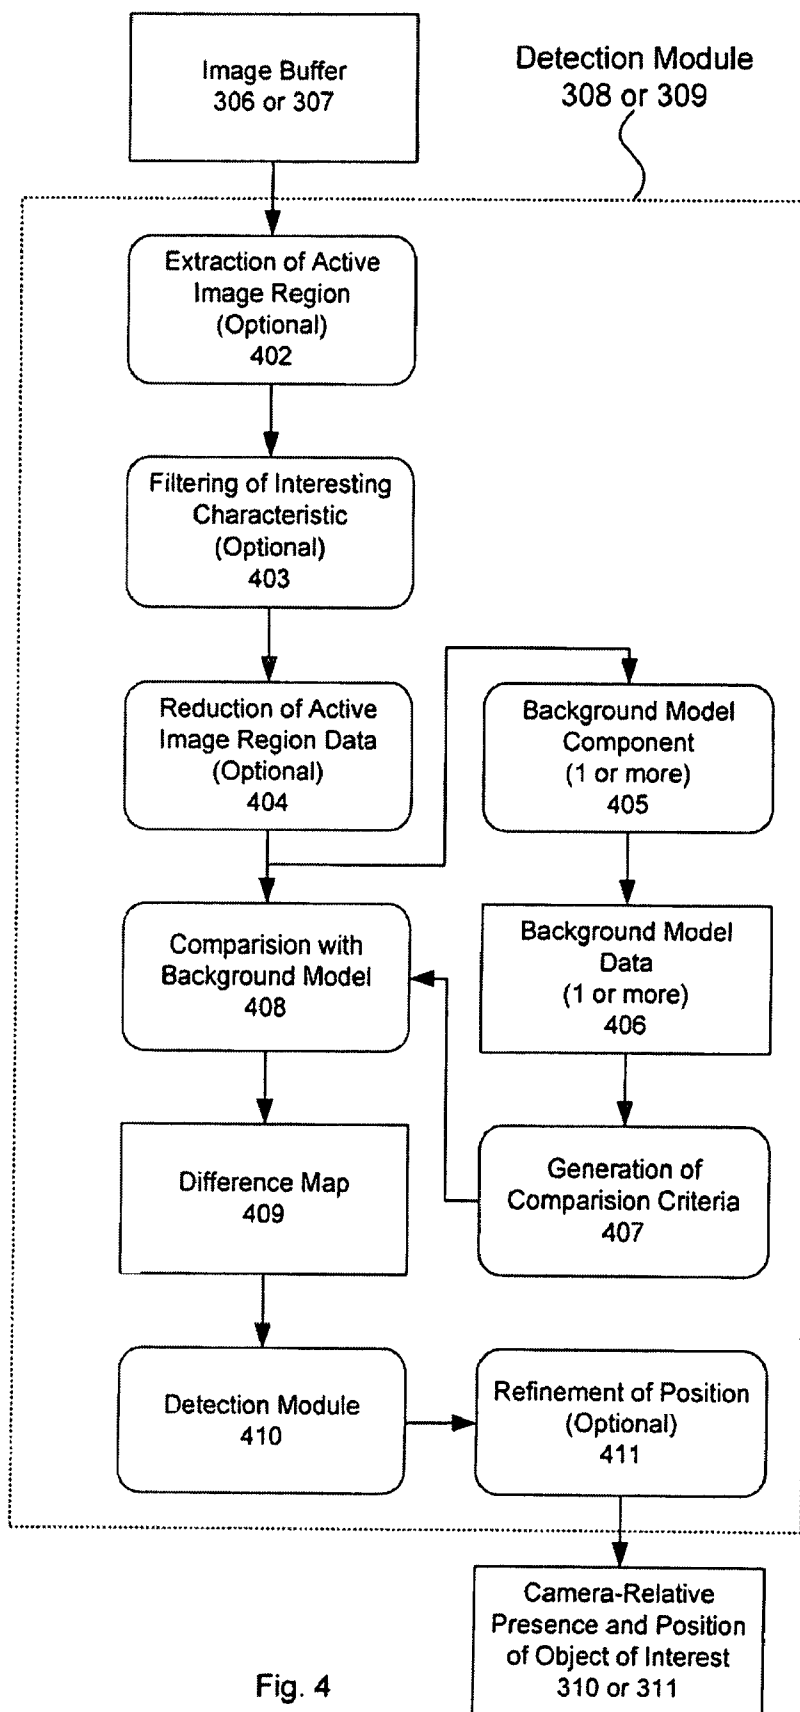


Fig. 4



Raw Image
Fig. 5A



Filtered Image
Fig. 5B



Low-Value Threshold
Fig. 5C



High-Value Threshold
Fig. 5D



Binary Difference
Fig. 5E

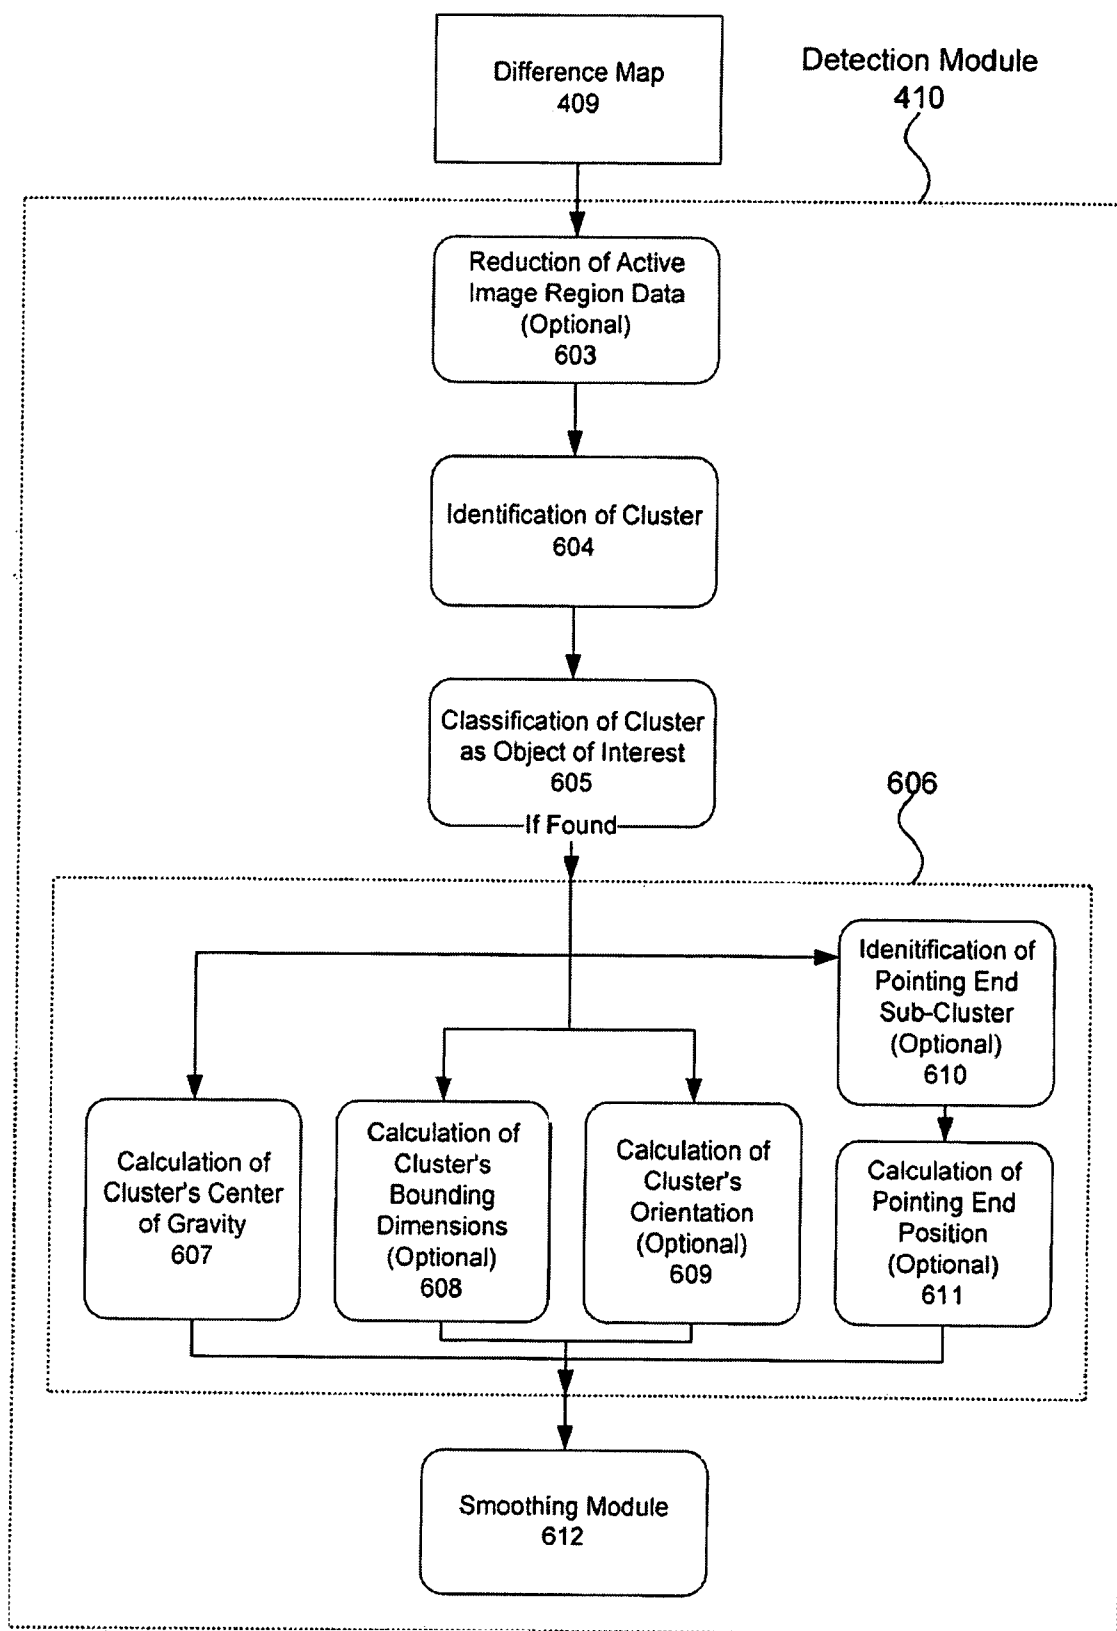
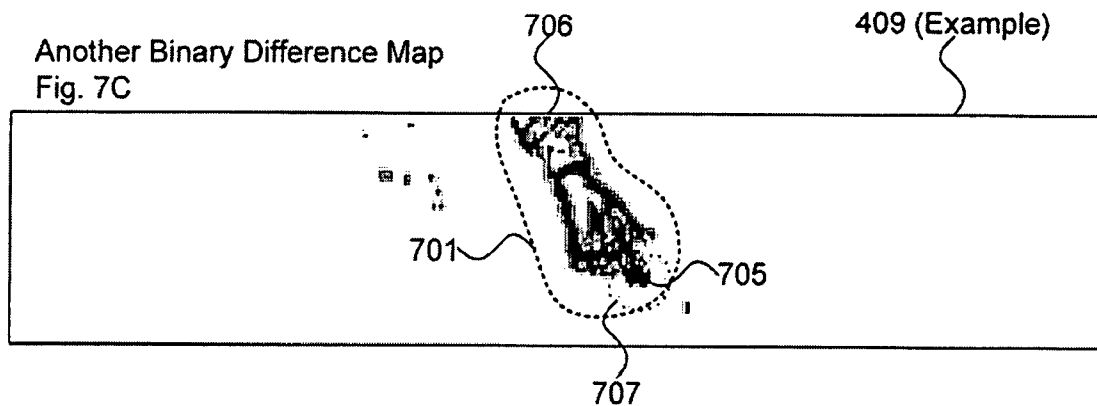




Fig. 6



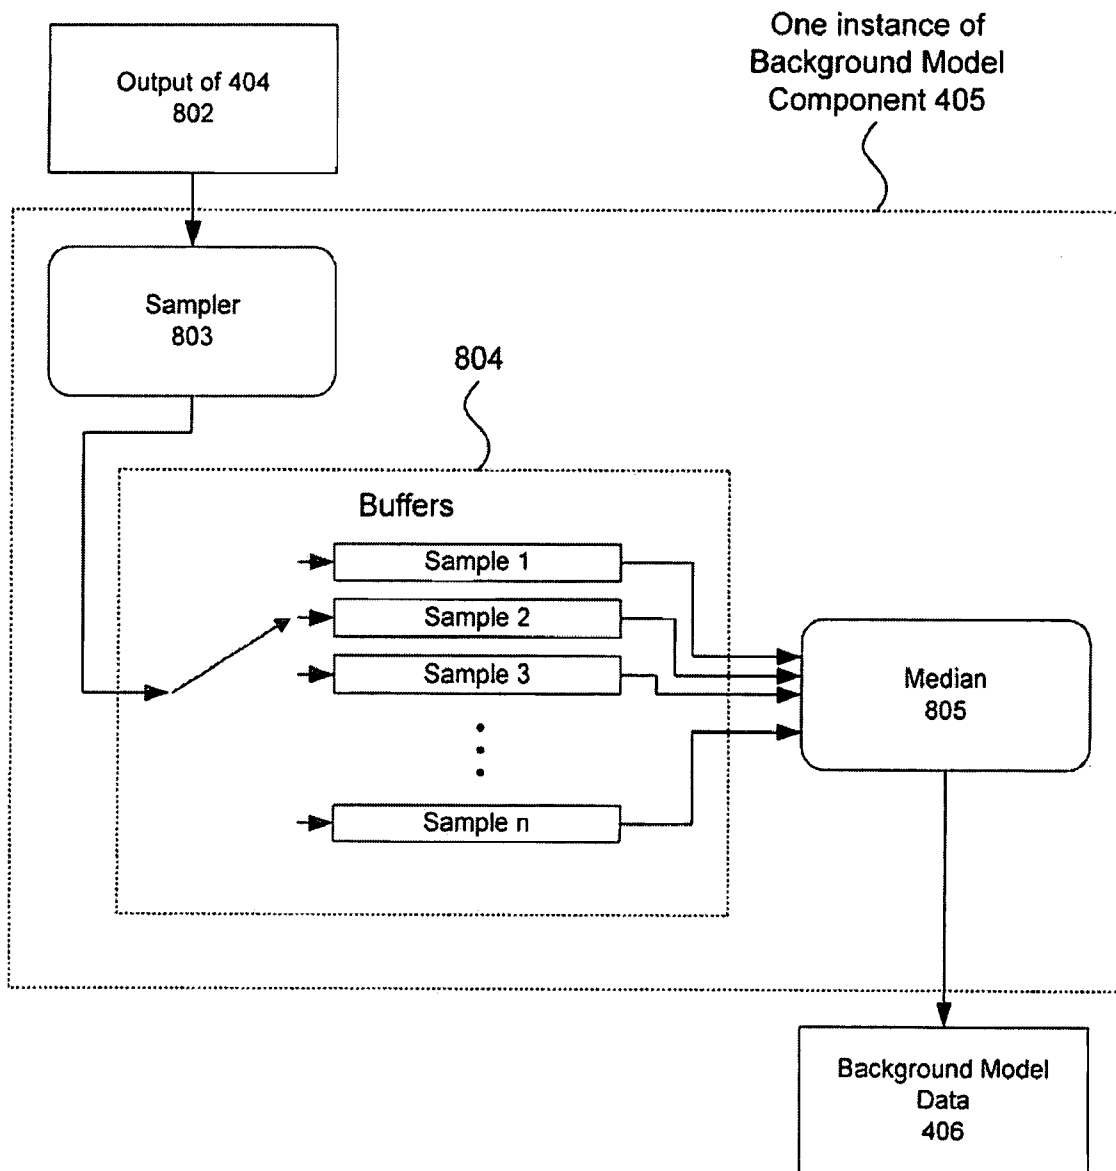


Fig. 8




Fig. 9A




Fig. 9B






Fig. 10






Fig. 12A




Fig. 12B




Fig. 13A




Fig. 13B




Fig. 13C




Fig. 13D

1

MULTIPLE CAMERA CONTROL SYSTEM**CROSS REFERENCE TO RELATED APPLICATION**

This application is a continuation of U.S. application Ser. No. 11/304,000, filed Dec. 19, 2005, which is a continuation of U.S. application Ser. No. 09/962,612, filed Sep. 26, 2001, which claims the benefit of U.S. Provisional Application No. 60/237,187, filed Oct. 3, 2000. All of these applications are incorporated by reference in their entirety.

TECHNICAL FIELD

This invention relates to an object tracking system, and more particularly to a video camera based object tracking and interface control system.

BACKGROUND

A variety of operating systems are currently available for interacting with and controlling a computer system. Many of these operating systems use standardized interface functions based on commonly accepted graphical user interface (GUI) functions and control techniques. As a result, different computer platforms and user applications can be easily controlled by a user who is relatively unfamiliar with the platform and/or application, as the functions and control techniques are generally common from one GUI to another.

One commonly accepted control technique is the use of a mouse or trackball style pointing device to move a cursor over screen objects. An action, such as clicking (single or double) on the object, executes a GUI function. However, for someone who is unfamiliar with operating a computer mouse, selecting GUI functions may present a challenge that prevents them from interfacing with the computer system. There also exist situations where it becomes impractical to provide access to a computer mouse or trackball, such as in front of a department store display window on a city street, or while standing in front of a large presentation screen to lecture before a group of people.

SUMMARY

In one general aspect, a method of tracking an object of interest is disclosed. The method includes acquiring a first image and a second image representing different viewpoints of the object of interest, and processing the first image into a first image data set and the second image into a second, image data set. The method further includes processing the first image data set and the second image data set to generate a background data set associated with a background, and generating a first difference map by determining differences between the first image data set and the background data set, and a second difference map by determining differences between the second image data set and the background data set. The method also includes detecting a first relative position of the object of interest in the first difference map and a second relative position of the object of interest in the second difference map, and producing an absolute position of the object of interest from the first and second relative positions of the object of interest.

The step of processing the first image into the first image data set and the second image into the second image data set may include determining an active image region for each of the first and second images, and extracting an active image data set from the first and second images contained within the

2

active image region. The step of extracting the active image data set may include one or more techniques of cropping the first and second images, rotating the first and second images, or shearing the first and second images.

In one implementation, the step of extracting the active image data set may include arranging the active image data set into an image pixel array having rows and columns. The step of extracting further may include identifying the maximum pixel value within each column of the image pixel array, and generating data sets having one row wherein the identified maximum pixel value for each column represents that column.

Processing the first image into a first image data set and the second image into a second image data set also may include filtering the first and second images. Filtering may include extracting the edges in the first and second images. Filtering further may include processing the first image data set and the second image data set to emphasize differences between the first image data set and the background data set, and to emphasize differences between the second image data set and the background data set.

Processing the first image data set and the second image data set to generate the background data set may include generating a first set of one or more background data sets associated with the first image data set, and generating a second set of one or more background data sets associated with the second image data set.

Generating the first set of one or more background data sets may include generating a first background set representing a maximum value of data within the first image data set representative of the background, and generating the second set of one or more background data sets includes generating a second background set representing a maximum value of data within the second image data set representative of the background. Generating further may include, for the first and second background sets representing the maximum value of data representative of the background, increasing the values contained within the first and second background sets by a predetermined value.

Generating the first set of one or more background data sets may include generating a first background set representing a minimum value of data within the first image data set representative of the background, and generating the second set of one or more background data sets may include generating a second background set representing a minimum value of data within the second image data set representative of the background. Generating further may include, for the first and second background sets representing the minimum value of data representative of the background, decreasing the values contained within the first and second background sets by a predetermined value.

Generating the first set of background data sets may include sampling the first image data set, and generating the second set of background data sets may include sampling the second image data set. Sampling may occur automatically at predefined time intervals, where each sample may include data that is not associated with the background.

Generating the first set of one or more background data sets may include maintaining multiple samples of the first image data set within each background data set, and generating the second set of one or more background data sets may include maintaining multiple samples of the second image data set within each background data set.

Generating each first background data set may include selecting from the multiple samples one value that is representative of the background for each element within the first image data set, and generating each second background data

set may include selecting from the multiple samples one value that is representative of the background for each, element within the second image data set. Selecting may include selecting the median value from all sample values in each of the background data sets.

In other implementations, generating may include comparing the first image data set to a subset of the background data set, and comparing the second image data set to a subset of the background data set.

In other implementations generating a first difference map further may include representing each element in the first image data set as one of two states, and generating a second difference map further may include representing each element in the second image data set as one of two states, where the two states represent whether the value is consistent with the background.

In still other implementations, detecting may include identifying a cluster in each of the first and second difference maps, where each cluster has elements whose state within its associated difference map indicates that the elements are inconsistent with the background.

Identifying the cluster further may include reducing the difference map to one row by counting the elements within a column that are inconsistent with the background. Identifying the cluster further may include identifying the column as being within the cluster and classifying nearby columns as being within the cluster. Identifying the column as being within the cluster also may; include identifying the median column.

Identifying the cluster further may include identifying a position associated with the cluster. Identifying the position associated with the cluster may include calculating the weighted mean of elements within the cluster.

Detecting further may include classifying the cluster as the object of interest. Classifying the cluster further may include counting the elements within the cluster and classifying the cluster as the object of interest only if that count exceeds a predefined threshold. Classifying the cluster further may include counting the elements within the cluster and counting a total number of elements classified as inconsistent within the background within the difference map, and classifying the cluster as the object of interest only if the ratio of the count of elements within the cluster over the total number of elements exceeds a predefined threshold.

The step of detecting further may include identifying a sub-cluster within the cluster that represents a pointing end of the object of interest and identifying a position of the sub-cluster.

In the above implementations, the object of interest may be a user's hand, and the method may include controlling, an application program using the absolute position of the object of interest.

The above implementations further may include acquiring a third image and a fourth image representing different viewpoints of the object of interest, processing the third image into a third image data set and the fourth image into a fourth image data set, and processing the third image data set and the fourth image data set to generate the background data set associated with the background. The method also may include generating a third difference map by determining differences between the third image data set (and the background data set, and a fourth difference map by determining differences between the fourth image data set and the background data set, and detecting a third relative position of the object of interest in the third difference map and a fourth relative position of the object of interest in the fourth difference map. The

absolute position of the object of interest may be produced from the first, second, third and fourth relative positions of the object of interest.

As part of this implementation, the object of interest may be a user's hand, and also may include controlling an application program using the absolute position of the object of interest.

In another aspect, a method of tracking an object of interest controlled by a user to interface with a computer is disclosed. The method includes acquiring images from at least two viewpoints, processing the acquired images to produce an image data set for each acquired image, and comparing each image data set to one or more background data sets to produce a difference map for each acquired image. The method also includes detecting a relative position of an object of interest within each difference map, producing an absolute position of the object of interest from the relative positions of the object of interest, and using the absolute position to allow the user to interact with a computer application.

Additionally, this method may include mapping the absolute position of the object of interest to screen coordinates associated with the computer application, and using the mapped position to interface with the computer application. This method also may include recognizing a gesture associated with the object of interest by analyzing changes in the absolute position of the object of interest, and combining the absolute position and the gesture to interface with the computer application.

In another aspect, a multiple camera tracking system for interfacing with an application program running on a computer is disclosed. The multiple camera tracking system includes two or more video cameras arranged to provide different viewpoints of a region of interest and are operable to produce a series of video images. A processor is operable to receive the series of video images and detect objects appearing in the region of interest. The processor executes a process to generate a background data set from the video images, generate an image data set for each received video image and compare each image data set to the background data set to produce a difference map for each image data set, detect a relative position of an object of interest within each difference map, and produce an absolute position of the object of interest from the relative positions of the object of interest and map the absolute position to a position indicator associated with the application program.

In the above implementation, the object of interest may be a human hand. Additionally, the region of interest may be defined to be in front of a video display associated with the computer. The processor may be operable to map the absolute position of the object of interest to the position indicator such that the location of the position indicator on the video display is aligned with the object of interest.

The region of interest may be defined to be any distance in front of a video display associated with the computer, and the processor may be operable to map the absolute position of the object of interest to the position indicator such that the location of the position indicator on the video display is aligned to a position pointed to by the object of interest. Alternatively, the region of interest may be defined to be any distance in front of a video display associated with the computer, and the processor may be operable to map the absolute position of the object of interest to the position indicator such that movements of the object of interest are scaled to larger movements of the location of the position indicator on the video display.

The processor may be configured to emulate a computer mouse function. This may include configuring the processor to emulate controlling buttons of a computer mouse using

5

gestures derived from the motion of the object of interest. A sustained position of the object of interest for a predetermined time period may trigger a selection action within the application program.

The processor may be configured to emulate controlling buttons of a computer mouse based on a sustained position of the object of interest for a predetermined time period. Sustaining a position of the object of interest within the bounds of an interactive display region for a predetermined time period may trigger a selection action within the application program.

The processor may be configured to emulate controlling buttons of a computer mouse based on a sustained position of the position indicator within the bounds of an interactive display region for a predetermined time period.

In the above aspects, the background data set may include data points representing at least a portion of a stationary structure. In this implementation, at least a portion of the stationary structure may include a patterned surface that is visible to the video cameras. The stationary structure may be a window frame. Alternatively, the stationary structure may include a strip of light.

In another aspect, a multiple camera tracking system for interfacing with an application program running on a computer is disclosed. The system includes two or more video cameras arranged to provide different viewpoints of a region of interest and are operable to produce a series of video images. A processor is operable to receive the series of video images and detect objects appearing in the region of interest. The processor executes a process to generate a background data set from the video images, generate an image data set for each received video image, compare each image data set to the background data set to produce a difference map for each image data set, detect a relative position of an object of interest within each difference map, produce an absolute position of the object of interest from the relative positions of the object of interest, define sub regions within the region of interest, identify a sub region occupied by the object of interest, associate an action with the identified sub region that is activated when the object of interest occupies the identified sub region, and apply the action to interface with the application program.

In the above implementation, the object of interest may be a human hand. Additionally, the action associated with the identified sub region may emulate the activation of keys of a keyboard associated with the application program. In a related implementation, sustaining a position of the object of interest in any sub region for a predetermined time period may trigger the action.

The details of one or more implementations are set forth in the accompanying drawings and the description below. Other features and advantages will be apparent from the description and drawings, and from the claims.

DESCRIPTION OF DRAWINGS

FIG. 1 shows the hardware components of a typical implementation of the multicamera control system, and their typical physical layout.

FIG. 2A shows the typical geometric relationship between the cameras and various image regions of FIG. 1.

FIG. 2B shows a typical image captured by one of the cameras of FIG. 1.

FIG. 3 is a low diagram showing the processes that are performed, typically within a microcomputer program associated with the multicamera control system.

FIG. 4 is a flow diagram showing a portion of the process shown in FIG. 3 in greater detail, and in particular, the pro-

6

cesses involved in detecting an object and extracting its position from the image signals captured by the cameras.

FIG. 5A shows sample image data, presented as a gray-scale bitmap image, acquired by a camera and generated by part of the process shown in FIG. 4.

FIG. 5B shows sample image data, presented as a gray-scale bitmap image, generated by part of the process shown in FIG. 4.

FIG. 5C shows sample image data, presented as a gray-scale bitmap image, generated by part of the process shown in FIG. 4.

FIG. 5D shows sample image data, presented as a gray-scale bitmap image, generated by part of the process shown in FIG. 4.

FIG. 5E shows sample data, presented as a binary bitmap image, identifying those pixels that likely belong to the object that is being tracked in the sample, generated by part of the process shown in FIG. 4.

FIG. 6 is a flow diagram showing a portion of the process described in FIG. 4 in greater detail, and in particular, the processes involved in classifying and identifying the object given a map of pixels that have been identified as likely to belong to the object that is being tracked, for example given the data shown in FIG. 5E.

FIG. 7A shows the sample data presented in FIG. 5E, presented as a binary bitmap image, with the identification of those data samples that the processes shown in FIG. 6 have selected as belonging to the object in this sample.

FIG. 7B shows the sample data presented in FIG. 5E, presented as a bar graph, with the identification of those data samples that the processes outlined in FIG. 6 have selected as belonging to the object, with specific points in the graph being identified.

FIG. 7C shows a difference set of sample data, presented as a binary bitmap image, with the identification of those data samples that the processes shown in FIG. 6 have selected as belonging to the object and key parts of the object in this sample.

FIG. 8 is a flow diagram that shows a part of the process shown in FIG. 4 in greater detail, and in particular, the processes involved in generating and maintaining a description of the background region over which the object occludes.

FIG. 9A shows the geometry on which Eq. 3 is based, that is, an angle defining the position of the object within the camera's field of view, given the location on the image plane where the object has been sensed.

FIG. 9B shows the geometry on which Eq. 4, 5 and 6 are based, that is, the relationship between the positions of the cameras and the object that is being tracked.

FIG. 10 is a graph illustrating Eq. 8, that is, the amount of dampening that may be applied to coordinates given the change in position of the object to refine the positions.

FIG. 11A is an example of an application program that is controlled by the system, where the object of interest controls a screen pointer in two dimensions.

FIG. 11B shows the mapping between real-world coordinates and screen coordinates used by the application program in FIG. 11A.

FIGS. 12A and 12B are examples of an application program that is controlled by the multicamera control system, where the object of interest controls a screen pointer in a three dimensional virtual reality environment.

FIG. 13A shows the division of the region of interest into detection planes used by a gesture detection method to identify a gesture that may be associated with the intention to activate.

FIG. 13B shows the division of the region of interest into detection boxes used by a gesture detection method to identify a gesture that may be associated with selecting a cursor direction.

FIG. 13C shows an alternate division of the region of interest into direction detection boxes used by a gesture detection method to identify a gesture that may be associated with selecting a cursor direction.

FIG. 13D illustrates in greater detail the relationship of neighboring divisions of FIG. 13C.

Like reference symbols in the various drawings indicate like elements.

DETAILED DESCRIPTION

FIG. 1 shows a multicamera motion tracking and control system 100 interfaced with an image viewing system. In this implementation two cameras 101 and 102 scan a region of interest 103. A controlled or known background 104 surrounds the region of interest 103. An object of interest 105 is tracked by the system when it enters the region of interest 103. The object of interest 105 may be any generic object inserted into the region of interest 103, and is typically a hand or finger of a system user. The object of interest 105 also may be a selection device such as a pointer.

The series of video images acquired from the cameras 101 and 102 are conveyed to a computing device or image processor 106. In this implementation, the computing device is a general-purpose computer that runs additional software that provides feedback to the user on a video display 107.

FIG. 2A illustrates a typical implementation of the multicamera control system 100. The two cameras 101 and 102 are positioned outside of the region of interest 103. The cameras are oriented so that the intersection 204 of their field of views (205 for camera 101, 206 for camera 102) completely encompasses the region of interest 103. The orientation is such that the cameras 101, 102 are rotated on axes: that are approximately parallel. In this example, a floor or window ledge and sidewalls provide a controlled background 104 having distinct edges. The corresponding view captured by camera 101 is shown in FIG. 2B. While not shown, it should be understood that the view captured by camera 102 is a mirror image of the view captured by camera 101. The controlled background 104 may not cover the camera's entire field of view 205. For each camera, an active image region 208 is found that is entirely contained within the controlled background 104, and also contains the entire region of interest 103. The background 104 is controlled so that a characteristic of the background can be modeled, and the object of interest 105, either in part or in whole, differs from the background 104 in that characteristic. When the object, of interest 105 appears within the region of interest 103, the object 105 will occlude a portion of the controlled background 104 within the active image region 208 of each camera 101, 102. In the location of the occlusion, either as a whole or in parts, the captured images will, in terms of the selected characteristic, be inconsistent with the model of the controlled background 104.

In summary, the object of interest 105 is identified and, if found, its position within the active image region 208 of both cameras is calculated. Using the position data of each camera 101, 102, as well as the positions of the cameras relative to the region of interest 103, and parameters describing the cameras, the position of the object of interest 105 within the region of interest 103 is calculated.

The processes performed by the image processor 106 (FIG. 1), which may be implemented through a software process, or alternatively through hardware, are generally shown in FIG.

3. The camera images are simultaneously conveyed from the cameras 101, 102 and captured by image acquisition modules 304, 305 (respectively) into image buffers 306, 307 (respectively) within the image processor 106. Image detection modules 308, 309 independently detect the object of interest 105 in each image, and determine its position relative to the camera view. The relative position information 310, 311 from both camera views is combined by a combination module 312 and optionally refined, by a position refinement module 313, to determine at block 314, the global presence and position of the object of interest 105 within the region of interest 103. Optionally, specific gestures performed by the user may be detected in a gesture detection module 315. The results of the gesture detection process are then conveyed to another process or application 316, either on the same image processor 106 or to another processing device. The process of gesture detection is described in greater detail below.

Image detection modules 308 and 309 are identical in the processes that they execute. An implementation of these image detection modules 308, 309 is shown in FIG. 4. In block 402, the image processor 106 extracts, from the captured image data stored in the image buffers 306 or 307, the image data that corresponds to the active image region 208 (of FIG. 2B). The image may be filtered in a filtering process 403 to emphasize or extract the aspects or characteristics of the image where the background 104 and object of interest 105 differ, but are otherwise invariant within the background 104 over time. In some implementations, the data representing the active image region may also be reduced by a scaling module 404 in order to reduce the amount of computations required in later processing steps. Using the resulting data, the background 104 is modeled by one or more instances of a background model process at block 405 to produce one or more descriptions represented as background model data 406 of the controlled background 104. Therefore the background 104 is modeled in terms of the desired aspects or characteristics of the image. The background model(s) 406 are converted into a set of criteria in process 407. In a comparison process 408, the filtered (from process 403) and/or reduced (from module 404) image data is compared to those criteria (from process 407), and the locations where the current data is inconsistent with the background model data 406, that is where the criteria is not satisfied, are stored in an image or difference map 409. In detection module 410, the difference map 409 is analyzed to determine if any such inconsistencies qualify as a possible indication of an object of interest 105 and, if these criteria are satisfied, its position within the camera view (205 or 206) is determined. The position of the object 105 may be further refined (optionally) at block 411, which produces a camera-relative presence and position output 310 or 311 associated with the object of interest 105 (as described above with respect to FIG. 3).

In block 402 of FIG. 4, image processor 106 extracts the image data that corresponds to the active image region 208 (of FIG. 2B). The image data may be extracted by cropping, shearing, rotating, or otherwise transforming the captured image data. Cropping extracts only the portion of the overall image that is within the active image region 208. Bounds are defined, and any pixels inside the bounds are copied, unmodified, to a new buffer, while pixels outside of the bounds are ignored. The active image region 208 may be of arbitrary shape. Shearing and rotation reorder the data into an order that is more convenient for further processing, such as a rectangular shape so that it may be addressed in terms of rows and columns of pixels.

Rotation causes the contents of an image to appear as if the image has been rotated. Rotation reorders the position of pixels from (x,y) to (x',y') according to the following equation:

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} \cos\theta & -\sin\theta & 0 \\ \sin\theta & \cos\theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

where θ is the angle that the image is to be rotated.

If the cameras **101** and **102** are correctly mounted with respect to the region of interest **103**, the desired angle of rotation will typically be small. If the desired angle of rotation is small, shearing may be used to provide an approximation that is computationally simpler than rotation. Shearing distorts the shape of an image such that the transformed shape appears as if the rows and columns have been caused to slide over and under each other. Shearing reorders the position of pixels according to the following equations:

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & sh_x & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \text{ and } \begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 \\ sh_y & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

where sh_x represents the amount of horizontal shear within the image, and sh_y represents the amount of vertical shear within the image.

An implementation of the multicamera control system **100** applies in scenarios where the object of interest **105**, either in whole or in part, is likely to have either higher or lower luminance than the controlled background **104**. For example, the background **104** may be illuminated to create this scenario. A filtering block **403** passes through the luminance information associated with the image data. A single background model **406** represents the expected luminance of the background **104**. In practice, the luminance of the controlled background **104** may vary within the active image region **208**, therefore the background model **406** may store the value of the expected luminance for every pixel within the active image region **208**. The comparison criteria generation process **407** accounts for signal noise (above that which may be accounted for within the background model) and minor variability of the luminance of the controlled background **104** by modifying each luminance value from the background model **406**, thus producing the minimal luminance value that may be classified as being consistent with the background model **406**. For example, if the luminance of the controlled background **104** is higher than the luminance of the object of interest **105**, then processes block **407** decreases the luminance value of each pixel by an amount greater than the expected magnitude of signal noise and variability of luminance.

In some implementations of system **100**, the region of interest **103** is sufficiently narrow such that it may be modeled as a region of a plane. The orientation of that plane is parallel to the front and rear faces of the dotted cube that represents the region of interest **103** in FIG. 1. The active image region **208** may be reduced to a single row of pixels in the optional scaling module **404** if two conditions are satisfied: 1) the object of interest **105**, when it is to be detected, will occlude the background **104** in all rows of some columns of the active image region **208**, and 2) a single set of values in the background model **406** sufficiently characterizes an entire

column of pixels in the active image region **208**. The first condition is usually satisfied if the active image region **208** is thinner than the object of interest **105**. The second condition is satisfied by the implementation of blocks **403**, **405**, **406** and **407** described above. Application of the scaling module **404** reduces the complexity of processing that is required to be performed in later processes, as well as reducing the storage requirements of the background model(s) **406**.

The particular implementation of the scaling module **404** depends on the specifics of processing blocks **403**, **405**, **406** and **407**. If the luminance of the controlled background **104** is expected to be higher than that of the object of interest **105**, as described above, one implementation of the scaling module **404** is to represent each column by the luminance of greatest magnitude within that column. That is to say, for each column, the highest value in that column is copied to a new array. This process has the added benefit that the high-luminance part of the controlled background **104** need not fill the entire controlled background **104**.

An alternative implementation applies in scenarios where the controlled background **104** is static that is, contains no motion, but is not otherwise limited in luminance. A sample source image is included in FIG. 5A as an example. In this case, the object of interest, as sensed by the camera, may contain, or be close in magnitude to, the luminance values that are also found within the controlled background **104**. In practice, the variability of luminance of the controlled background **104** (for example, caused by a user moving in front of the apparatus thereby blocking some ambient light) may be significant in magnitude relative to the difference between the controlled background **104** and the object of interest **105**. Therefore, a specific type of filter may be applied in the filtering process **403** that produces results that are invariant to or de-emphasize variability in global luminance, while emphasizing parts of the object of interest **105**. A 3x3 Prewitt filter is typically used in the filtering process **403**. FIG. 5B shows the result of this 3x3 Prewitt filter on the image in FIG. 5A. In this implementation, two background models **406** may be maintained, one representing each of the high and low values, and together representing the range of values expected for each filtered pixel. The comparison criteria generation process **407** then decreases the low-value and increases the high-value by an amount greater than the expected magnitude of signal noise and variability of luminance. The result is a set of criterion, an example of which, for the low-value, is shown in FIG. 5C, and an example of which, for the high-value, is shown in FIG. 5D. These modified images are passed to the comparison process **408**, which classifies pixels as being inconsistent to the controlled background **104** if their value is either lower than the low-value criterion (FIG. 5C) or higher than the high-value criterion (FIG. 5D). The result is a binary difference map **409**, of which example corresponding to FIG. 5B is shown in FIG. 5E.

The preceding implementation allows the use of many existing surfaces, walls or window frames, for example, as the controlled background **104** where those surfaces may have arbitrary luminance, textures, edges, or even a light strip secured to the surface of the controlled background **104**. The above implementation also allows the use of a controlled background **104** that contains a predetermined pattern or texture, a stripe for example, where the above processes detect the lack of the pattern in the area where the object of interest **105** occludes the controlled background **104**.

The difference map **409** stores the positions of all pixels that are found to be inconsistent with the background **104** by the above methods. In this implementation, the difference map **409** may be represented as a binary image, where each

11

pixel may be in one of two states. Those pixels that are inconsistent with the background **104** are identified or “tagged” by setting the pixel in the corresponding row and column of the difference map to one of those states. Otherwise, the corresponding pixel is set to the other state.

An implementation of the detection module **410**, which detects an object of interest **105** in the difference map **409**, shown in FIG. 6. Another scaling module at block **603** provides an additional opportunity to reduce the data to a single dimensional array of data, and may optionally be applied to scenarios where the orientation of the object of interest **105** does not have a significant effect on the overall bounds of the object of interest **105** within the difference map **409**. In practice, this applies to many scenarios where the number of rows is less than or similar to the typical number of columns that the object of interest **105** occupies. When applied, the scaling module at block **603** reduces the difference map **409** into a map of one row, that is, a single dimensional array of values. In this implementation, the scaling module **603** may count the number of tagged pixels in each column of the difference map **409**. As an example, the difference map **409** of FIG. 7A is reduced in this manner and depicted as a graph **709** in FIG. 7B. Applying this optional processing step reduces the processing requirements and simplifies some of the calculations that follow.

Continuing with this implementation of the detection module **410**, it is observed that the pixels tagged in the difference map (**409** in example FIG. 7A) that are associated with the object of interest **105** will generally form a cluster **701**, however the cluster is not necessarily connected. A cluster identification process **604** classifies pixels (or, if the scaling module **603** has been applied, classifies columns) as to whether they are members of the cluster **701**. A variety of methods of finding clusters of samples exist and may be applied, and the following methods have been selected on the basis of processing simplicity. It is noted that, when the object of interest **105** is present, it is likely that the count of correctly tagged pixels will exceed the number of false-positives. Therefore, the median position is expected to fall somewhere within the object of interest **105**. Part of this implementation of the cluster identification process **604**, when applied to a map of one row (for example, where the scaling module at block **603** or **404** has been applied), is to calculate the median column **702** and tag columns as part of the cluster **701** (FIG. 7B) if they are within a predetermined distance **703** that corresponds to the maximum number of columns expected to be occupied. Part of this implementation of the cluster identification process **604**, when applied to a map of multiple rows, is to add tagged pixels to the cluster **703** if they meet a neighbor-distance criterion.

In this implementation, a set of criteria is received by a cluster classification process **605** and is then imposed onto the cluster **701** to verify that the cluster has qualities consistent with those expected of the object of interest **105**. Thus, process **605** determines whether the cluster **701** should be classified as belonging to the object of interest **105**. Part of this implementation of the cluster classification process **605** is to calculate a count of the tagged pixels within the cluster **701** and to calculate a count of all tagged pixels. The count within the cluster **701** is compared to a threshold, eliminating false matches in clusters having too few tagged pixels to be considered as an object of interest **105**. Also, the ratio of the count of pixels within the cluster **701** relative to the total count is compared to a threshold, further reducing false matches.

If the cluster **701** passes these criteria, a description of the cluster is refined in process block **606** by calculating the center of gravity associated with the cluster **701** in process

12

607. Although the median position found by the scaling module **603** is likely to be within the bounds defining the object of interest **105**, it is not necessarily at the object’s center. The weighted mean **710**, or center of gravity, provides a better measure of the cluster’s position and is optionally calculated within process **606**, as sub-process **607**. The weighted mean **710** is calculated by the following equation:

$$\bar{x} = \frac{\sum_{x=0}^{c-1} x \cdot C[x]}{\sum_{x=0}^{c-1} C[x]}$$

where: \bar{x} is the mean

c is the number of columns

C[x] is the count of tagged pixels in column x.

The cluster’s bounds **704** may also be optionally calculated within process **606**, shown as process **608**. The cluster **703** may include some false-positive outliers, so as part of this implementation, the bounds may be defined as those that encompass a predetermined percentile of the tagged pixels, or, in scenarios where relatively few pixels are expected to be tagged, encompasses those tagged pixels (or columns, if scaling module **603** is applied) that form tight sub-clusters, that is those tagged pixels (or columns) that have neighbors that are also tagged.

In addition to the middle and bound coordinates, the orientation of the object of interest **105** may optionally be inferred by calculation of the moments of the cluster. This calculation is represented by a cluster orientation calculation process at sub-process **609** within process **606**.

In some applications of the system **100**, the object of interest **105** is used as a pointer. In this case, the “pointing end” of the object **105** is desired and may also be determined by a pointing end calculation sub-process within process **606** if the region of interest **103** contains a sufficient number of rows and the number of rows has not been reduced. An example is depicted in FIG. 7C. The object of interest **105** will typically enter, or be constrained to enter, the active image region **208** from a known border of that region. The pointing end **705** (for example the user’s fingertip) of the object of interest **105** is likely to be the portion of the cluster **701** that is furthest from the region of entry **706** into the active image region **208**. The cluster **701** may include some false-positive outliers. As such, the pointing end **705** may be defined as the region **707** within the cluster **701** that encompasses multiple tagged pixels near the furthest bounding side of the cluster **701**, or, in scenarios where relatively few pixels are expected to be tagged, encompasses the furthest tagged pixels that form a tight sub-cluster; that is those tagged pixels that have neighbors that are also tagged. This sub-cluster is identified by a sub-cluster pointing end process **610**, and the position of the sub-cluster is found in process **611**.

Continuing with this implementation, a process implemented by a smoothing module **612** may optionally be applied to any or all of the positions found in process **606**. Smoothing is a process of combining the results with those solved previously so they move in a steady manner from frame to frame. The weighted mean coordinate **710**, found by the center of gravity determination process **607**, is dependent on many samples and therefore is inherently steady. The bound **704**, found by the cluster bounding dimension determination process **608**, and pointing end **705**, found by **611**, coordinates are dependent on relatively fewer members of the

cluster, and the state of a single pixel may have a significant effect. Since the size of the region occupied by the object of interest **105** is expected to remain relatively steady, smoothing may be applied to the distance between the bounds **704** measured relative to the cluster's weighted mean coordinate **710**. Since the shape and orientation of the object of interest **105** is expected to change less rapidly than the overall position object of interest **105**, smoothing may be applied to the distance of the pointing end **705** measured relative to the cluster's weighted mean coordinate **710**.

A process used in the center of gravity process **607** is Eq. 1 as follows:

$$s(t) = (\alpha \times r(t)) + ((1 - \alpha) \times s(t-1))$$

In Eq. 1, the smoothed value at time t ($s(t)$) is equal to one minus the scalar value (α) multiplied by the smoothed value at time minus one ($t-1$). This amount is added to the raw value at time t ($r(t)$) multiplied by a scalar (α) that is between zero and one.

Referring to FIG. 8, implementations of system **100** make use of, as described above, one or more background models **406** (FIG. 4). An implementation of the background model process or component **405** that generates the background model data **406** is shown in FIG. 8. This implementation of the background model component **405** automatically generates and dynamically updates the background model, allowing unattended operation of the system.

Input data **802** is provided by the output of scaling module **404** for this implementation of the background model component **405**. Input is available every frame, and is sampled in a sampling process **803**. The sample may contain the object of interest **105** occluding part of the controlled background **104**. For each pixel, a range of values may be a better representative of the background **104** than a single value. By including the effects of this range in the background model, the expansion in process **407** may be made tighter. Contributing multiple frames of data to the sample allows this range to be observed, but also increases the portion of the background **104** that is occluded by the object of interest **105** if the object of interest **105** is in motion while the frames are being sampled. The optimal number of frames to use is dependent on the expected motion of the object of interest **105** in the particular application of the system. In practice, for systems that are tracking a hand, 10 frames, representing approximately 0.33 seconds, is sufficient to observe the majority of that range without allowing motion of the object of interest to occlude an undue portion of the background. If the particular background model is to be compared in comparison process **408** as the upper bound on values that are considered to be consistent with the background **104**, then the maximum value of each pixel observed in the multiple frames may be recorded as the sample value. If the particular background model **406** is to be compared in process **408** as the lower bound on values that are considered to be consistent with the background **104**, then the minimum value of each pixel observed in the multiple frames may be recorded as the sample value.

In this implementation of the background model component **405**, samples from the sampling process **803** are added to a buffer **804** having storage locations to store n samples, where the oldest sample in the history is replaced. The history therefore contains n sampled values for each pixel. The span of time, d , represented in the buffer is dependent on the rate that new samples are acquired and added to the history, r , by Eq. 2, described as follows:

$$d = \frac{n}{r}$$

In this implementation, a median process block **805** selects, for each pixel, a value that it determines is representative of the controlled background **104** at the location represented by that pixel. One method of selecting a value representative of the controlled background **104** within process block **805** is to select the median value of the n samples of each pixel. For any pixel, a number of then sampled values in the buffer **804** may represent the object of interest **105**. Duration d is selected so that it is: unlikely that the object of interest **105** will occlude any one pixel of the controlled background **104** for an accumulated duration of $d/2$ or longer within any time-span of d . Therefore, for any pixel, the majority of the sampled values will be representative of the background **104**, and therefore the median of the sampled values will be a value representative of the background **104**.

The background model component **405** is adaptive, and any changes to the background **104** will be reflected in the output of median process block **805** once they have been observed for time of $d/2$. This system does not require that the entire controlled background **104** be visible when initialized, the object of interest **105** may be present when initialized, however it does require that samples be observed for time of d before providing output. Optionally, the constraint may be applied that the object of interest **105** must be absent when the system is initialized, in which case the first observed sample values may be copied into all n samples of the buffer **804**, allowing the system to produce an output sooner.

The duration that any one pixel of the controlled background **104** will be occluded by the object of interest **105**, and therefore the duration, is dependent on the particular application of the system. The number of samples, n , can be scaled for the memory buffer and processing power available.

The preceding discussion presents one implementation of obtaining the position of the object of interest **105** within and relative to the images acquired by the cameras **101** and **102**. If the object of interest **105** was successfully detected and its coordinates found in both cameras views **205** and **206** by detection modules **308** and **309** of FIG. 3, then the combination of these coordinates is sufficient to recover the position of the object of interest **105** within the region of interest **103**. In the implementation outlined in FIG. 3, the position of the object of interest **105** is calculated in combination module **312**.

Turning to FIGS. 9A and 9B, an implementation of the combination module **312** is shown. For each camera **101** and **102**, the position p **902** of the object of interest **105** on the camera's image plane **904** is converted to an angle **905**, which is referred in this description as beta (β), and is measured on the reference plane whose normal is defined by the axes of the rotations of the cameras **101**, **102**. (In practice, the axes are not precisely parallel and do not exactly define a single plane, however the process described herein is tolerant of that error). By approximating the camera **101**, **102** as an ideal pinhole model of the camera, that angle (β), relative to the vector **906** defining the orientation of the camera, is approximated.

15

Eq. 3, as shown in FIG. 9A, illustrates an approximation calculation as follows:

$$\beta = \tan^{-1}\left(\frac{f}{p}\right)$$

To approximate the angle beta (β), the inverse tangent is applied to the quantity of the focal length (f) divided by the position p on the image plane projected onto the intersection of the reference plane and the image plane.

For maximum precision, the intrinsic camera parameters (location of the principal point and scale of image) and radial distortion caused by the lens should be corrected for by converting the distorted position (as represented by the relative position information 310, 311) to the ideal position. More specifically, the ideal position is the position on the image plane 904 that the object 105 would be projected if the camera 101, 102 had the properties of an ideal pinhole camera, whereby Eq. 3 will produce the exact angle. One set of correction equations are presented in Z. Zhang, A Flexible New Technique for Camera Calibration, Microsoft Research, <http://research.microsoft.com/~zhang>, which is incorporated by reference. For many applications of the system, the approximation has been found to provide sufficient precision without this correction noted above.

Continuing with the description of combination module 312, a reference vector 907, as illustrated in FIG. 9B, is defined such that it passes through the positions of both cameras 101 and 102 on the reference plane where the reference plane is defined such that the axis of rotation of the cameras define the normal of the reference plane. The angles 908 that the cameras are rotated are measured relative to the reference vector 907.

A formula for measurement of the angles is shown in Eq. 4:

$$\alpha = \beta_0 + \beta$$

Measurement of the angle alpha (α) is equal to the angle beta_not (β_0) and the angle beta (β).

Eq. 4 is applied to measure the angles 909 of the object of interest 105 relative to the reference vector 907. That angle is referred to by the alpha (α) symbol herein. The angle alpha 909 for each camera 101 and 102, and the length of the reference vector 907, are sufficient to find the position of the object of interest 105 on the reference plane, by Eq. 5 and Eq. 6.

Eq. 5 calculates the offset of the object of interest (y) by the formula.

$$y = \frac{w \tan \alpha_A \tan \alpha_B}{\tan \alpha_A + \tan \alpha_B}$$

The offset (y) is equal to the reciprocal of the tangent of the angle (α_A) for camera A 101 and the tangent of the angle (α_B) for camera B 102 multiplied by the vector length 907 (w), the tangent of the angle (α_A) for camera A 101 and the tangent of the angle (α_B) for camera B 102.

16

Eq. 6 calculates the offset of the object of interest (x_A) as follows:

$$x_A = \frac{y}{\tan \alpha_A}$$

In Eq. 6, the offset (x_A) is measured by the offset from Eq. 5 (y) divided by the tangent of the angle (α_A) for camera A 101.

The position of the object 105 on the axis perpendicular to, the reference plane may be found by Eq. 7, which is applied to the position in each image, using the distance of the object of interest 105 from the camera.

$$z = l \frac{p}{f}$$

In Eq. 7, the position (z) is calculated as the position (p) on the image plane projected onto the vector of the image plane perpendicular to, that use in Eq. 3 divided by the focal length (f) multiplied by the distance of the object of interest 105 from the camera (l).

These relations provide a coordinate of the object of interest 105 relative to Camera A 101. Knowing the position and size of the region of interest 103 relative to Camera A 101, the coordinate may be converted so that it is relative to the region of interest 103, 312 of FIG. 3.

Smoothing may optionally be applied to these coordinates in refinement module 313 of the implementation of this system shown in FIG. 3. Smoothing is a process of combining the results with those solved previously so that motion is steady from frame to frame. One method of smoothing for these particular coordinate values (x_A , y, z found by combination module 312) is described herein. Each of the components of the coordinate values associated with the object of interest 105 that is x, y, and z, are smoothed independently and dynamically. The degree of dampening S is calculated by Eq. 8, where S is dynamically and automatically adjusted in response to the change in position is calculated as follows:

$$S = \begin{cases} S_A & \text{if } (D \leq D_A) \\ \alpha S_B + (1 - \alpha) S_A & \text{where } \alpha = \frac{D - D_A}{D_B - D_A} \text{ if } (D_A < D < D_B) \\ S_B & \text{if } (D \geq D_B) \end{cases}$$

$$D = |r(t) - s(t - 1)|$$

In Eq. 8, s(t) is the smoothed value at time t, r(t) is the raw value at time t, D_A and D_B are thresholds, and S_A and S_B define degrees of dampening.

Two distance thresholds, D_A and D_B , as shown in FIG. 10, define three ranges of motion. A change in position that is less than D_A , motion is heavily dampened 1001 by S_A , thereby reducing the tendency of a value to switch back and forth between two nearby values (a side effect of the discrete sampling of the images). A change in position greater than D_B is lightly dampened 1002 by S_B , or not dampened. This reduces or eliminates lag and vagueness that is introduced in some other smoothing procedures. The degree of dampening is varied for motion between D_A and D_B , the region marked as 1003, so that the transition between light and heavy dampening is less noticeable. The scalar a, which is applied to Eq. 1, is found by Eq. 9 as follows:

17

$$a = \frac{e(1-S)}{S}$$

In Eq. 9, scalar (α) is bound such that equal to or greater than zero, and less than or equal to one, the dampening value of S is found by Eq. 8, and e is the elapsed time since the previous frame.

These coordinates 314 of the object of interest 105, if found, are typically conveyed to another process such as a user application program 316 for use. They may be conveyed to another process executing on the same image processor 106 as the above calculations where performed, or to another computing device. The method in which the data are conveyed to the application program 316 may include emulation of a traditional user input device (including mouse and keyboard), allowing the system to provide control of existing control functions within the application program 316. The coordinates 314 of the object of interest 105 may be calculated for every video frame captured by the cameras, where one video frame is typically captured 30 times or more every second. This results in little latency between the user's actions and the application's reactions.

In a typical implementation of the system, the application program 316 provides user feedback by displaying to the video display 107 a visual representation of an indicator. The indicator is caused to move such that its position and motion mimics the motion of the object of interest 105 (typically the user's hand).

In one variation of this form of user interface, the indicator, such as a mouse pointer, is shown in front of other graphics, and its movements are mapped to the two dimensional space defined by the surface of the screen. This form of control is analogous: to that provided by a computer mouse, such as that used with the Microsoft® Windows® operating system. An example feedback image of an application that uses this style of control is shown as 1102 in FIG. 11A.

Referring to FIG. 11A (and briefly to FIG. 3), the image processor 106 also includes an optional coordinate re-mapping process 317 (FIG. 3). The coordinate re-mapping process 317 is operable to remap the global presence and position coordinates 314 (associated with the object of interest 105) into the position where the indicator 1101 (such as a cursor or mouse pointer) is overlaid onto the image 1102 by way of Eq. 10 for the x coordinate, and the equivalent of this equation for the y coordinate, as follows:

$$x_c = \begin{cases} 0 & \text{if } x_h < b_l \\ \frac{x_h - b_l}{b_r - b_l} & \text{if } b_l \leq x_h \leq b_r \\ 1 & \text{if } x_h > b_r \end{cases}$$

In Eq. 10, x_h is the coordinate position 314 associated with the object 105, x_c is the cursor position on the screen, mapped 0-1, and b_l and b_r are the positions of the left and right bounds of a sub-region within the region of interest 103. As illustrated in FIG. 11B, the entire region of the display 1102 is represented by a sub-region 1103 contained entirely within the region of interest 103. Positions (for example, position A 1105) within the sub-region 1103 are linearly mapped to positions (for example, 1106) within the display 1102. Positions (for example position B 1107) outside the sub-region 1103 but still within the region of interest 103 are mapped to

18

the nearest position (for example, 1108) on the border of the display region 1102. This reduces the likelihood of the user unintentionally removing the object of interest 105 (usually the user's hand or pointing finger) from the sub-region while attempting to move the indicator 1101 to a position near a border of the display.

In scenarios where the region of interest 103 is immediately in front of the video display 107, the sub-region 1103 may be defined to be aligned to the video display 107, so that the indicator 1101 will appear to be aligned with the object of interest 105. If the region of interest 103 is relatively thin, for example less than 5 cm, and the sub-region 1103 is defined in this way, then the system approximates, in terms of user-interaction, a "touch-screen" without limitations on the size of the video display 107, and without requiring direct contact between the user and video display's 107 surface (for example, the video display and user may be on opposite sides of a window). As will be appreciated, the system 100 can be used with a variety of video display sizes, and may include not only computer monitors (whether CRT or LCD type displays), but also may include rear projection style television monitors, large flat screen LCD monitors, and forward projection style presentation systems.

In scenarios where the region of interest 103 is not immediately in front of a large video display 107, and the active image region 208 is sufficiently deep that the orientation of the object of interest is found in the orientation calculation process 609, a vector may be extended from the object of interest's position to the video display 167 using the angle of orientation to detect the position on the video display that the user is "pointing to."

Most often, however, the active image region 208 is not sufficiently deep to accurately calculate the orientation in process block 609. In these scenarios, where the region of interest 103 is not immediately in front of a large video display 107 and the orientation is not calculated Eq. 0 may be applied where the sub-region 1103 is smaller than the video display. The processor then maps the absolute position of the object of interest 105 to the position indicator such that movements of the object of interest 105 are scaled to larger movements of the location of the position indicator on the video display, which allows the entire area of the video display to be easily reached by the user (for example the sub region 1103 may be defined to be at most 750 mm in width and proportional in height, a size that is easily reached by most users). When setup in this way, the system still provides the user the feeling of "pointing to the screen."

In another variation of this form of user interface, the user causes a representation of an indicator to move within a representation of a three dimensional virtual environment (examples are presented in FIG. 12A and FIG. 12B). The virtual environment may be rendered using projective transforms, so that the depths of the virtual environment are implied by the image presented on the video display 107. Techniques for rendering this sort of virtual environment include OpenGL. Eq. 10 is used to remap the x, y, and z coordinates (the sub-region 1103 becomes, for example, a cube).

Applications that are controlled by a movable on screen indicator (for example, FIGS. 11A, 12A, and 12B), whose control has been discussed, typically present graphic representations of data or interactive elements (for example, a button 1109 or an object representation 1202). The user is expected to cause the indicator 1101 to be positioned over one of these objects, or if a three-dimensional virtual environment is presented, touches or interacts, with the object. For a two-dimensional interface, this condition may be detected by

19

comparing the remapped indicator position **1106** to the bounds (for example, **1110**) of the graphic representation of the object, where this condition is true if the indicator position is within the object bounds. For the three-dimensional interface, this condition may be detected by comparing the bounds **1203** of either the entire indicator **1101**, or if finer control is required, a part of the indicator, with the bounds **1204** of the object **1202**. The user optionally receives feedback indicating that the cursor is positioned over an object. Feedback may be of a variety of forms, including an audio cue and/or a change in the graphical representation of either or both the cursor and object. The user may then activate, manipulate, or move the object that is under the cursor. The user is expected to indicate his intention to activate, manipulate, or move the object by performing a gesture.

The motion of the object of interest **105** may optionally be interpreted and classified by the gesture detection module **315** as described above with respect to FIG. 3. The gesture detection process **315** may utilize the data produced from any component of the system. The final coordinates **314**, image coordinates **310** and **311**, or a combination of **310**, **311**, and **314**, may be sampled over time and provided as input to the gesture detection process **315**. A variety of gestures (for example, "hovering" and "poking") have been successfully detected using this data as input to a gesture detection process **315**.

In scenarios where the application's state (that is, whether of not the indicator **1101** is over a button **1109**) is known and is conveyed to the gesture detection module **315**. One gesture that the user performs to indicate the intention to activate the object (for example screen objects **1109**, **1202**) that is under the cursor **1101** is to cause the cursor to hover over the object (examples **1109**, **1202**) for longer than a predefined duration. This gesture performed by the user is detected by monitoring the application's state and triggering the gesture when the application state remains unchanged for the predetermined duration. The application need not be created specifically for the multicamera control system **100**, as techniques exist that can unobtrusively monitor an application's state (in the Windows operating system by setting a "hook" using the Windows SDK function "SetWindowsHookEx") and emulating a mouse "click" (in the Windows operating system by using the Windows SDK function "SendInput").

In some scenarios, the application state may not be available and may not be monitored. In this case, some exemplary gestures that indicate the intention to activate the object (for example screen objects **1109**, **1202**) under the cursor **1101** are holding the hand stationary ("hovering"), or poking the hand quickly forward and back.

A method by which "hovering" has been detected is by keeping a history of the position of the object of interest **105**, where that history contains all records of the position and state for a predefined duration of time, ending with the most recent sample. That duration represents the minimum duration that the user must hold the hand stationary. The minimum and maximum position, separately in each of the three (x,y,z) dimensions, is found within the history. If the object of interest **105** was present within the region of interest **103** in all samples of the history, and the distance between the minimum and maximum is within a predefined threshold for each of the three dimensions, then the "hovering" gesture is reported. Those distance thresholds represent the maximum amount that the object of interest **105** is allowed to move, plus the maximum amount of variation (or "jitter") expected to be introduced into the hand position by the various components of the system. The typical method in which this gesture is reported, where the system is emulating a mouse as described

20

above, is to emulate a mouse "click." Gestures representing additional operations of the mouse, "double clicks," and "dragging," have also been detected and those operations have been emulated.

In addition, gestures that are independent of the position of the indicator relative to an object may optionally be detected and given meaning by the application that may or may not be dependent on the application's state. An application that uses this style of interaction typically does not explicitly use or display the object of interest's position **317** or other positions. These applications can be wholly or primarily controlled with only the interpretations of the positions made by this system. These applications also need not be created specifically for this system because the interpretations made by this system can be used to simulate an action that would be performed on a traditional user input device, such as a keyboard or joystick.

Many useful interpretations depend directly on the absolute position of the object of interest **105** within the region of interest **103**. (Alternately, the indicator position **1105** within the sub-region **1103** may be used in an equivalent manner). One method of making these interpretations is to define boxes, planes, or other shapes. A state is triggered on if the position (for example the position defined by block **314**, or alternately by the remapped coordinates from remapping process **317**) of the object of interest **105** is found to be within a first box (or beyond the border defined by the first plane), and had not been in the immediately preceding observation (either because it was elsewhere within the region of interest **103**, or was not detected). This state is maintained until the hand position is not found to be within a second box (or beyond the border defined by the second plane), at which time the state is triggered off. The second box must contain the entire first box, and is typically larger. The use of a larger box reduces occurrences of the state unintentionally triggering on and off when the object of interest **105** is detected to be near the border of the boxes, where a very small motion or minor noise in the image signals would otherwise cause the position **317** to otherwise drift in and out of the box. Typically one of three methods of interpreting this state is used, depending on the intended use of the gesture. In one method, the gesture directly reflects the state with an on and off trigger. When emulating a keyboard key or joystick fire button, it is "pressed" when the state is triggered on, and "released" when the state is triggered off. In another method, the gesture is only triggered by the transition of the state from off to on. When emulating a keyboard key or joystick button, the key is "clicked." Although the duration and off state are not reported to the application, they are maintained so that the gesture will not be repeated until after the state is triggered off, so that each instance of the gesture requires a clearly defined intent by the user. A third method is to trigger the gesture when by the transition of the state from off to on, and to periodically re-trigger the gesture at predefined intervals so long as the state remains on. This emulates that way in which, holding a key down on a keyboard, causes the character to repeat in some applications.

One way in which boxes or planes, for the above techniques, may be defined within the region of interest **103** is as follows. By defining a first plane (**1501** in FIG. 13A) and second plane **1502** that divides the region of interest into "fire" **1503** and "neutral" **1504** regions (the gesture reported when the object of interest **105** is in the region **1505** between the planes depends on the previous positions of the object, as described above), the above technique can detect the object of interest **105** (typically a hand) "pushing" forward, which is one gesture for emulating a fire button on a joystick, or causing the application to respond in a way that is commonly

associated with the pressing of a joystick button (for example, the firing of a weapon in a video game).

Another technique in which boxes or planes, for the above techniques, may be defined within the region of interest **103** is as follows. Planes of the first type **1506**, **1507**, **1508**, **1509** are defined that separate each of the left, right, top and bottom portions of the region of interest **103**, overlapping in the corner regions as illustrated in FIG. **13B**. Planes of the second type are labeled as **1510**, **1511**, **1512**, **1513**. Each pair of first and second planes is processed independently. This combination of planes emulates the four directional cursor keys, where a hand in a corner triggers two keys, commonly interpreted by many applications as the four secondary 45 degree (diagonal) directions. Emulating the keyboard cursor in this method allows a variety of existing applications to be controlled by system **100**, including, for example, Microsoft® PowerPoint® which responds to the emulated cursor keys (e.g. the up and down arrow keys) by advancing to the next or previous slide in a presentation sequence.

Another method of emulating control of discreet directions applies for applications that expect the four 45 degree direction states to be explicitly represented. Boxes **1514**, **1515**, **1516**, **1517** are defined for each of the four primary (horizontal and vertical) directions, with additional boxes **1518**, **1519**, **1520**, **1521** defined for each of the secondary 45 degree (diagonal) directions as illustrated FIG. **13C**. For clarity, only boxes of the first type are illustrated. A gap is placed between these boxes. FIG. **13D** illustrates how neighboring boxes are defined. The gap between boxes of the first type **1522**, **1523** assures that the user intentionally causes the object of interest **105** to enter the box, while the gap **1524** is filled by overlapping boxes of the second type **1525**, **1526**, so that the system will report the previous gesture until the user was clearly intended to move the object of interest **105** into either a neighboring box or the central neutral region. This combination of buttons can be used to emulate an eight-directional joystick pad.

A wider class of gestures depend on motion instead of or in addition to position. An example is the gesture of “swiping the hand to the left.” This is a one gesture to convey to an application that it is to return to a previous page or state. Through emulation of a keyboard and mouse, this gesture may be used to control information presentation software, in particular Microsoft® PowerPoint®, to go to the previous slide of a presentation sequence. Through emulation of a keyboard and mouse, this gesture causes a web browser to perform the action associated with its “back” button. Similarly, the gesture of “swiping the hand to the right” is one gesture to convey to an application that the user desires to go to the next page or state. For example, this gesture causes presentation software to go to the next slide of a presentation sequence, and causes browser software to go to the next page.

One method for detecting “swiping the hand to the left” is as follows. A thin stripe along the leftmost part of the region of interest **103** is defined as the left-edge region. The position (for example the position defined by block **314**, or alternately by the remapped coordinates from remapping process **317**) of the object of interest **105** is represented as the following three states:

1. Object of interest is present and not in side the left-edge region
2. Object of interest is present and inside the left-edge region
3. Object of interest is not present within the hand detection region.

A transition from state 1 to state 2 above causes the gesture detection module **315** to enter a state whereby it starts a timer

and waits for the next transition. If a transition to state 3 is observed within a predetermined duration of time, the “swiping the hand off to the left” gesture is reported to have occurred. This technique is typically duplicated for the right, upper, and lower edges, and, because the hand position is found in three dimensions, also duplicated to detect “pulling the hand back.”

A variety of gesture detection techniques have been discussed. Still other gesture detection techniques (for example, Hidden Markov Layers) are described in research literature, and may be applied in the various implementations of the system **100** described herein.

Referring back to FIGS. **1** and **3**, another implementation of the multicamera control system **100** is described in further detail. While FIG. **1** shows a two camera system, it should be understood that the image processor **106** can be configured to receive input from more than two cameras, and may for particular applications include four (4) or more video cameras. In the four camera implementation, components **304-311** of FIG. **3** are duplicated to support the two additional cameras. Additionally, the combination module **312** is configured to receive four sets of camera-relative presence and position data (similar to data **310** and **311**) associated with the object of interest **105** being tracked. The techniques and equations (in particular, Eq. 5 and Eq. 6) previously described can be applied to the additional pair(s) of cameras, where the output of the combination module **312** is the average of all the position from each of the camera pairs. The gesture detection module **315** is similarly reconfigured to receive four sets of camera-relative presence and position data **310**, **311** from the two additional detection modules (similar to **308**, **309**) which are substantially similar to detection modules **310** and **311**.

The output from the image processor **106**, which now includes processed object position coordinates and gesture information associated with four cameras, can be used by another process or user application program **316**. The formulas and geometry (described above) used to calculate coordinate information associated with the object of interest **105** from the two additional cameras are also used.

In one implementation using four cameras, the two additional cameras are positioned at the bottom two corners within the controlled background **104** and are oriented such that the region of interest **103** is within the field of view **205** of each camera. The advantage of a four camera system is that the position of the object of interest **105** can be tracked with greater accuracy. Thus, the application program may include more screen objects with increased density on the video display **107** because the increased tracking accuracy allows objects that are close in proximity to be correctly selected by small movements with the object of interest **105**. Moreover, the two additional cameras reduce errors in tracking the object of interest **105** when a portion of the object of interest **105** is occluded within the field of view **205** associated with one or more of the other cameras.

While a number of implementations have been described, it will be understood that various modifications may be made. Accordingly, other implementations are within the scope of the following claims.

What is claimed is:

1. A computer-implemented method comprising:

accessing, using a processor, first and second images each capturing a different view of an object and a background;

detecting, using the processor, for each of the first and second images, a difference between a filtered luminance characteristic and a background model;

23

determining, using the processor, a position of the object based on the detected differences; and
controlling, using the processor, an application based on the determined position of the object.

2. The method of claim 1, wherein the background is a controlled background.

3. The method of claim 2, wherein the controlled background comprises a strip of light.

4. The method of claim 2, wherein the controlled background comprises a surface, wall or window frame having arbitrary luminance, textures or edges.

5. The method of claim 2, wherein the controlled background exhibits a predetermined pattern or texture.

6. The method of claim 5, wherein the predetermined pattern is a stripe pattern.

7. The method of claim 5, wherein detecting, for each of the first and second images, the occlusion of the background further comprises:

detecting, for each of the first and second images, a region that does not exhibit the predetermined pattern or texture; and

outputting, for each of the first and second images, data associated with the region as the occlusion.

8. The method of claim 1, wherein the first and second images are acquired using first and second cameras, respectively, and wherein the determined position of the object is a relative position with respect to the first and second cameras, respectively.

9. The method of claim 1, wherein the determined position of the objection is an absolute position.

10. The method of claim 1, further comprising detecting a gesture of the object, wherein the application is controlled further based on the gesture.

11. The method of claim 1, wherein detecting, for each of the first and second images, the difference between the filtered luminance characteristic and the background model further comprises filtering the first and second images to extract the luminance characteristic of the first and second images where the background and the object differ, but are otherwise invariant within the background over time.

12. The method of claim 1, further comprising processing the first and second images to exclude pixels outside bounds of the background.

13. The method of claim 1, wherein the first and second images each comprise a single row of pixels.

14. The method of claim 1, wherein detecting, for each of the first and second images, the difference between the luminance characteristic and the background model further comprises:

tagging a pixel of the first or second image to a first state if the pixel is consistent with the background model; and
tagging the pixel to a second state if the pixel is inconsistent with the background model.

15. The method of claim 1, further comprising providing, as the background, a strip of light that is capable of being detected in the first and second images.

16. An apparatus comprising:

one or more processing devices and a computer-readable medium coupled to the one or more processing devices, the medium storing instructions which, when executed by the one or more processing devices, cause the one or more computers to:

acquire first and second images each capturing a different view of an object and a background;

detect, for each of the first and second images, a difference between a filtered luminance characteristic and a background model; and

24

means for determining a position of the object based on the detected differences and controlling an application based on the determined position of the object.

17. A device comprising:

a camera configured to acquire first and second images each capturing a different view of an object and a background; and

a processor configured to:

detect, for each of the first and second images, a difference between a filtered luminance characteristic and a background model,

determine a position of the object based on the detected differences, and

control an application based on the determined position of the object.

18. A computer readable medium encoded with a computer program comprising instructions that, when executed, operate to cause a computer to:

acquire first and second images each capturing a different view of an object and a background;

detect, for each of the first and second images, a difference between a filtered luminance characteristic and a background model;

determine a position of the object based on the detected differences; and

control an application based on the determined position of the object.

19. A computer readable medium encoded with a computer program comprising instructions that, when executed, operate to cause a computer to:

acquire first and second images each capturing a different view of an object;

filter a luminance characteristic of each of the first and second images;

determine, based on filtering the luminance characteristic, a position of the object; and

control an application based on the determined position of the object.

20. A computer readable medium encoded with a computer program comprising instructions that, when executed, operate to cause a computer to:

acquire first and second images each capturing a different view of an object;

detect a difference between the first and second images;

determine a position of the object based on the difference; and

control an application based on the determined position of the object.

21. A computer readable medium encoded with a computer program comprising instructions that, when executed, operate to cause a computer to:

acquire first and second images each capturing a different view of an object and a background, the background being oriented parallel to a reference plane defined with respect to an imaging device configured to acquire at least one of the first and second images;

filter a luminance characteristic of each of the first and second images;

determine, based on the filtered luminance characteristic, a position of the object along an axis that is perpendicular to the reference plane; and

control an application based on the determined position of the object.

22. A computer readable medium encoded with a computer program comprising instructions that, when executed, operate to cause a computer to;

25

define first and second planes oriented parallel with a reference plane defined with respect to an imaging device configured to acquire at least one of the first and second images;

acquire first and second images each capturing a different view of an object and a background;

filter a luminance characteristic of each of the first and second images;

determine, based on the filtered luminance characteristic, a state of the object based on a position of the object along an axis that is perpendicular to the reference plane in relation to the first and second planes; and

generate first or second user inputs to an application if the object is determined to be in first or second states, respectively.

23. A computer readable medium encoded with a computer program comprising instructions that, when executed, operate to cause a computer to acquire an image of a view of an object;

generate a background model, further comprising:

adding the image to a buffer configured to store image data associated with the view;

selecting a value representative of a background for each portion of the view;

outputting, as the background model, the selected values;

determine a position of the object based on comparing the acquired image with the background model; and

control an application based on the determined position of the object.

24. A computer readable medium encoded with a computer program comprising instructions that, when executed, operate to cause a computer to;

acquire first through fourth images, the first and second images each capturing different views of an object at a first position, and the third and fourth images each capturing different views of the object at a second position;

filter a luminance characteristic of each of the first through fourth images;

determine, based on the filtered luminance characteristic, a motion of the object; and

control an application based on the determined motion.

25. A computer-implemented method of tracking an object of interest, the method comprising:

accessing a first image and a second image representing different viewpoints of the object of interest, the first and second images being accessed by a processor coupled to a computer, the processor further performing operations comprising:

processing the first image into a first image data set and the second image into a second image data set;

processing the first image data set and the second image data set to generate a background data set associated with a background;

generating a first difference map by determining differences between the first image data set and the background data set and a second difference map by determining differences between the second image data set and the background data set;

detecting a first relative position of the object of interest in the first difference map and a second relative position of the object of interest in the second difference map; and

producing an absolute position of the object of interest from the first and second relative positions of the object of interest.

26

26. A computer readable medium encoded with a computer program comprising instructions that, when executed, operate to cause a computer to;

acquire a first image capturing a view of an object and a background;

detect, for the first image, an occlusion of the background;

determine a position of the object based on the detected occlusion; and

control an application based on the determined position of the object and further based on a filtered luminance characteristic of the first image.

27. The method of claim 1, wherein accessing first and second images comprises acquiring the first and second images from at least one camera configured to transfer data to the processor.

28. The method of claim 25, wherein accessing first and second images comprises acquiring the first and second images from at least one camera configured to transfer data to the processor.

29. The medium of claim 21, further comprising instructions to: determine a second position of the object along the axis; and detect, based on the position and the second position, a hovering of the object with respect to the reference plane.

30. The medium of claim 22, wherein instructions to determine the state of the object comprises instructions to determine a gesture made by the object.

31. The medium of claim 30, wherein the gesture comprises a poking motion made by the object.

32. The medium of claim 23, wherein the state of the object is a motion of the object.

33. The medium of claim 24, wherein the motion is a gesture.

34. A computer readable medium encoded with a computer program comprising instructions that, when executed, operate to cause a computer to:

access a first image and a second image representing different viewpoints of the object of interest;

process the first image into a first image data set and the second image into a second image data set;

process the first image data set and the second image data set to generate a background data set associated with a background;

generate a first difference map by determining differences between the first image data set and the background data set and a second difference map by determining differences between the second image data set and the background data set;

detect a first relative position of the object of interest in the first difference map and a second relative position of the object of interest in the second difference map; and

produce an absolute position of the object of interest from the first and second relative positions of the object of interest.

35. A system comprising:

one or more computers; and

a computer-readable medium coupled to the one or more computers having instructions stored thereon which, when executed by the one or more computers, cause the one or more computers to perform operations comprising:

accessing a first image and a second image representing different viewpoints of the object of interest;

processing the first image into a first image data set and the second image into a second image data set;

27

processing the first image data set and the second image data set to generate a background data set associated with a background;

generating a first difference map by determining differences between the first image data set and the background data set and a second difference map by determining differences between the second image data set and the background data set;

5

28

detecting a first relative position of the object of interest in the first difference map and a second relative position of the object of interest in the second difference map; and

producing an absolute position of the object of interest from the first and second relative positions of the object of interest.

* * * * *

UNITED STATES PATENT AND TRADEMARK OFFICE
CERTIFICATE OF CORRECTION

PATENT NO. : 7,555,142 B2
APPLICATION NO. : 11/932869
DATED : June 30, 2009
INVENTOR(S) : Evan Hildreth

Page 1 of 3

It is certified that error appears in the above-identified patent and that said Letters Patent is hereby corrected as shown below:

Title page, item [57] under “Abstract”, line 10, delete “background-” and insert --background--

Claim 8

At Col. 23, lines 26-28, delete:

“8. The method of claim 1, wherein the first and second images are acquired using first and second cameras, respectively, and wherein the determined position of the object is a relative position with respect to the first and second cameras, respectively.”

and insert:

--8. The method of claim 1, wherein the first and second images are acquired using first and second cameras, respectively, and
 wherein the determined position of the object is a relative position with respect to the first and second cameras, respectively.--

Claim 9

At Col. 23, line 30, delete “objection” and insert --object--

Claim 18

At Col. 24, line 18, delete “to;” and insert --to:--

Claim 19

At Col. 24, line 30, delete “to;” and insert --to:--

Claim 20

At Col. 24, line 41, delete “to;” and insert --to:--

Claim 21

At Col. 24, line 52, delete “to;” and insert --to:--

Claim 22

At Col. 24, line 67, delete “to;” and insert --to:--

Claim 23

At Col. 25, lines 18-19, delete:

“23. A computer readable medium encoded with a computer program comprising instructions that, when executed, operate to cause a computer to acquire an image of a view of an object;
 generate a background model, further comprising:
 adding the image to a buffer configured to store image data associated with the view;
 selecting a value representative of a background for each portion of the view;
 outputting, as the background model, the selected values;
 determine a position of the object based on comparing the acquired image with the background model; and
 control an application based on the determined position of the object.”

and insert:

--23. A computer readable medium encoded with a computer program comprising instructions that, when executed, operate to:
 cause a computer to acquire an image of a view of an object;
 generate a background model, further comprising:
 adding the image to a buffer configured to store image data associated with the view;
 selecting a value representative of a background for each portion of the view;
 outputting, as the background model, the selected values;
 determine a position of the object based on comparing the acquired image with the background model; and
 control an application based on the determined position of the object.--

Claim 24

At Col. 25, line 34, delete “to;” and insert --to:--

Claim 25

At Col. 25, line 48, delete “b” and insert --by--

Claim 26

At Col. 26, line 3, delete “to;” and insert --to:--

Claim 29

At Col. 26, lines 21-24, delete:


“29. The medium of claim 21, further comprising instructions to: determine a second position of the object along the axis; and detect, based on the position and the second position, a hovering of the object with respect to the reference plane.”

and insert:

--29. The medium of claim 21, further comprising instructions to:
determine a second position of the object along the axis; and
detect, based on the position and the second position, a hovering of the object with
respect to the reference plane.--

Signed and Sealed this

Thirteenth Day of April, 2010

A handwritten signature in black ink that reads "David J. Kappos". The signature is written in a cursive, flowing style with a large initial 'D' and a stylized 'K'.

David J. Kappos
Director of the United States Patent and Trademark Office